

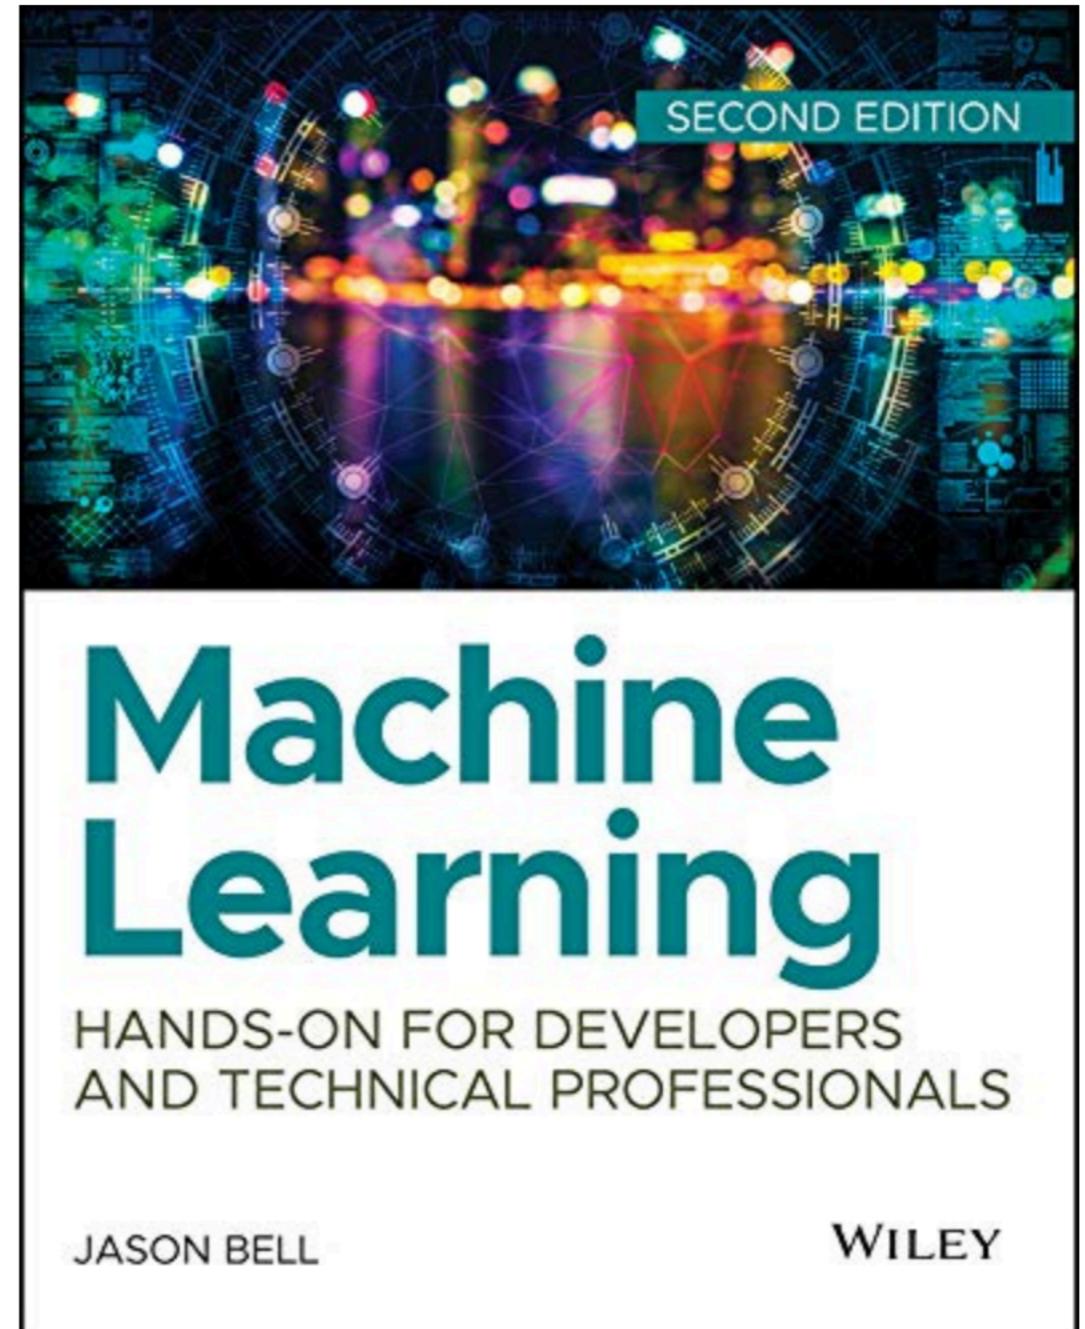
Hello Cleveland!

Work for

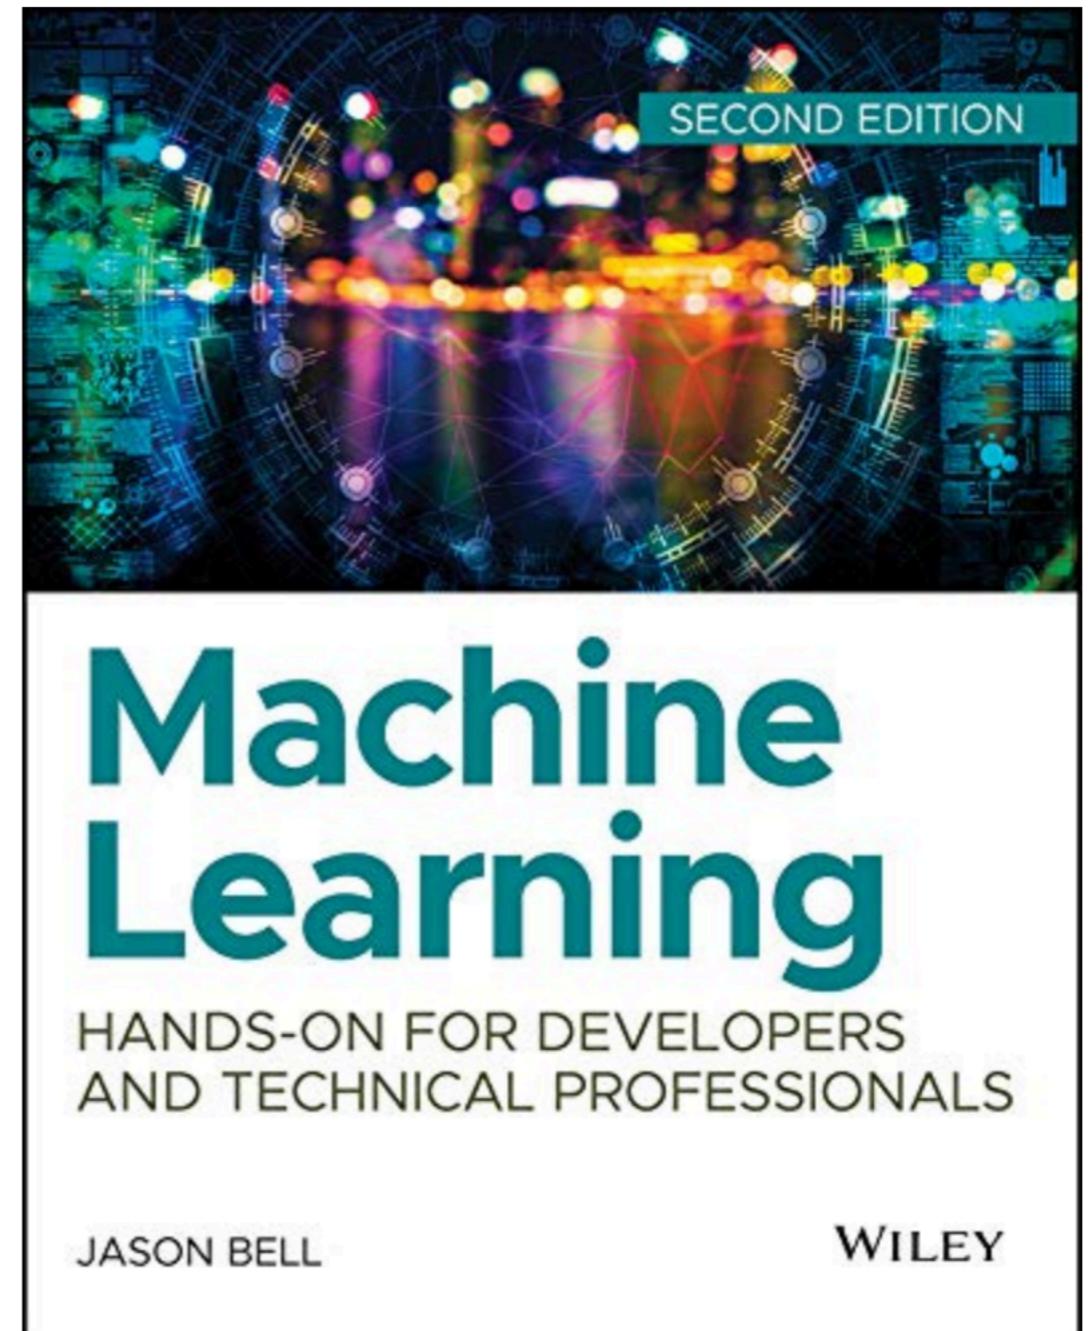


digitalis.io

Wrote this >>>



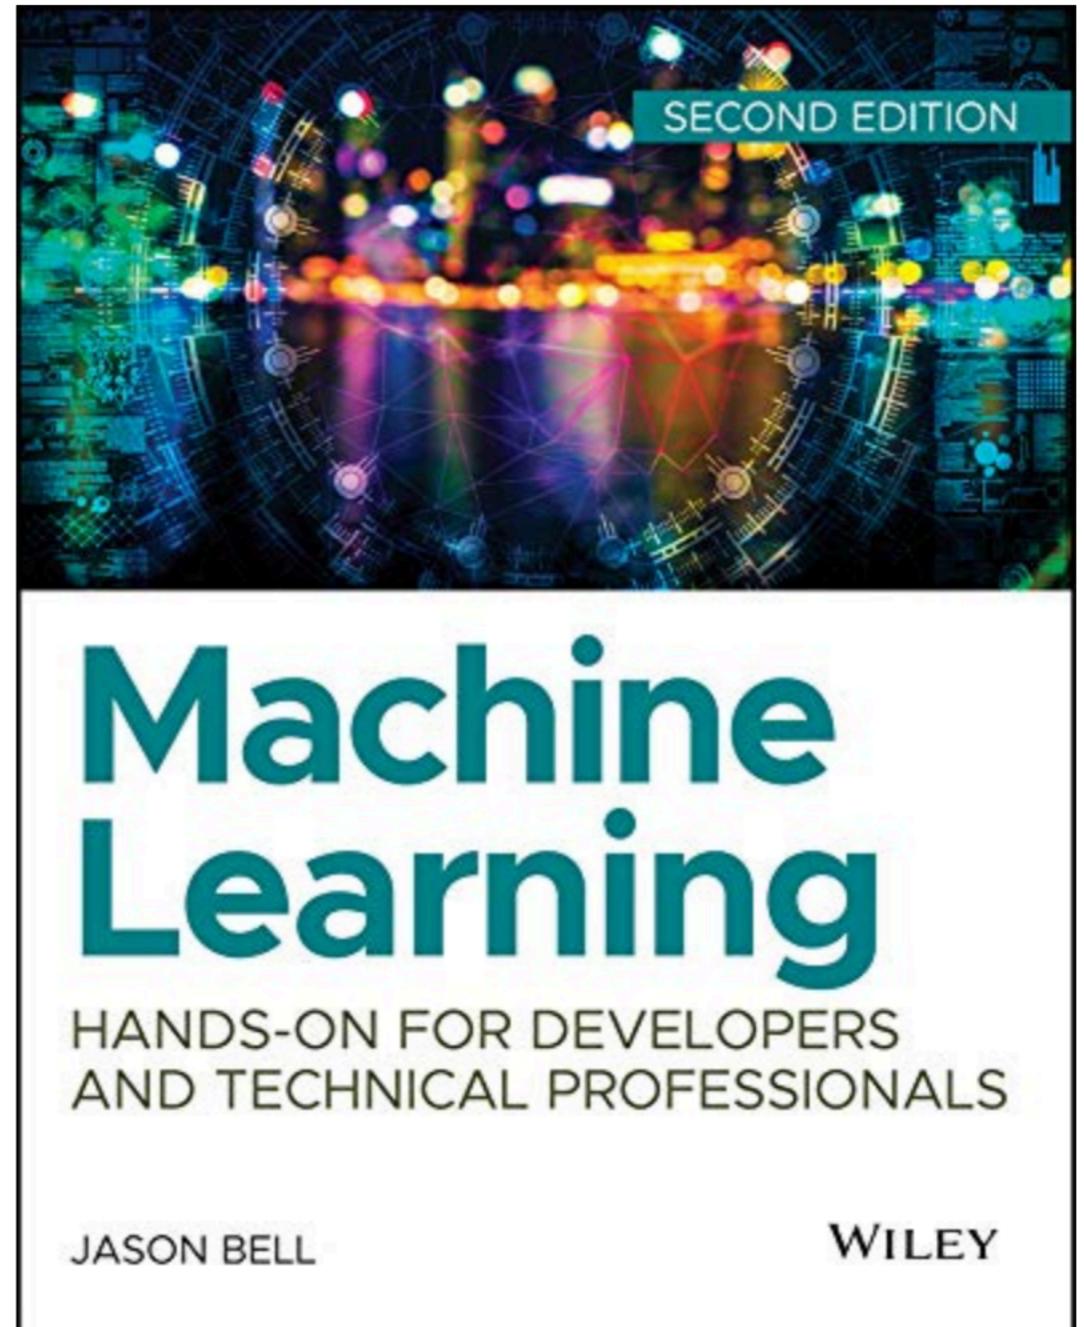
You can heckle on Twitter....
@jasonbelldata



The Plan

1. A quick overview of project and my original talk.
2. Look at how to make it better with Kafka's components.

Think of this talk as a conceptual how-to and thoughts about awkward data.



THIS IS A TRUE STORY.

**The events depicted took place in
London in 2016.**

**At the request of the client,
the names have been changed.**

**Out of respect for the data,
the rest has been told exactly
as it occurred.**

The Data

CSV Payloads are GZipped and sent to an endpoint via HTTP.

The data contains live flight searches and are batched together in chunks of a certain duration.

LHR,IST,2020-08-01,169.99,Y

**There may be 3 rows or 20,000 rows.
It's variable and we can't control that.**

I call this “The Fun Bit”

So this is what we did as a proof of concept.

And boy did I learn.



Thoughts, Experiences and
Considerations with Throughput
and Latency on High Volume
Stream Processing Systems Using
Cloud Based Infrastructures.

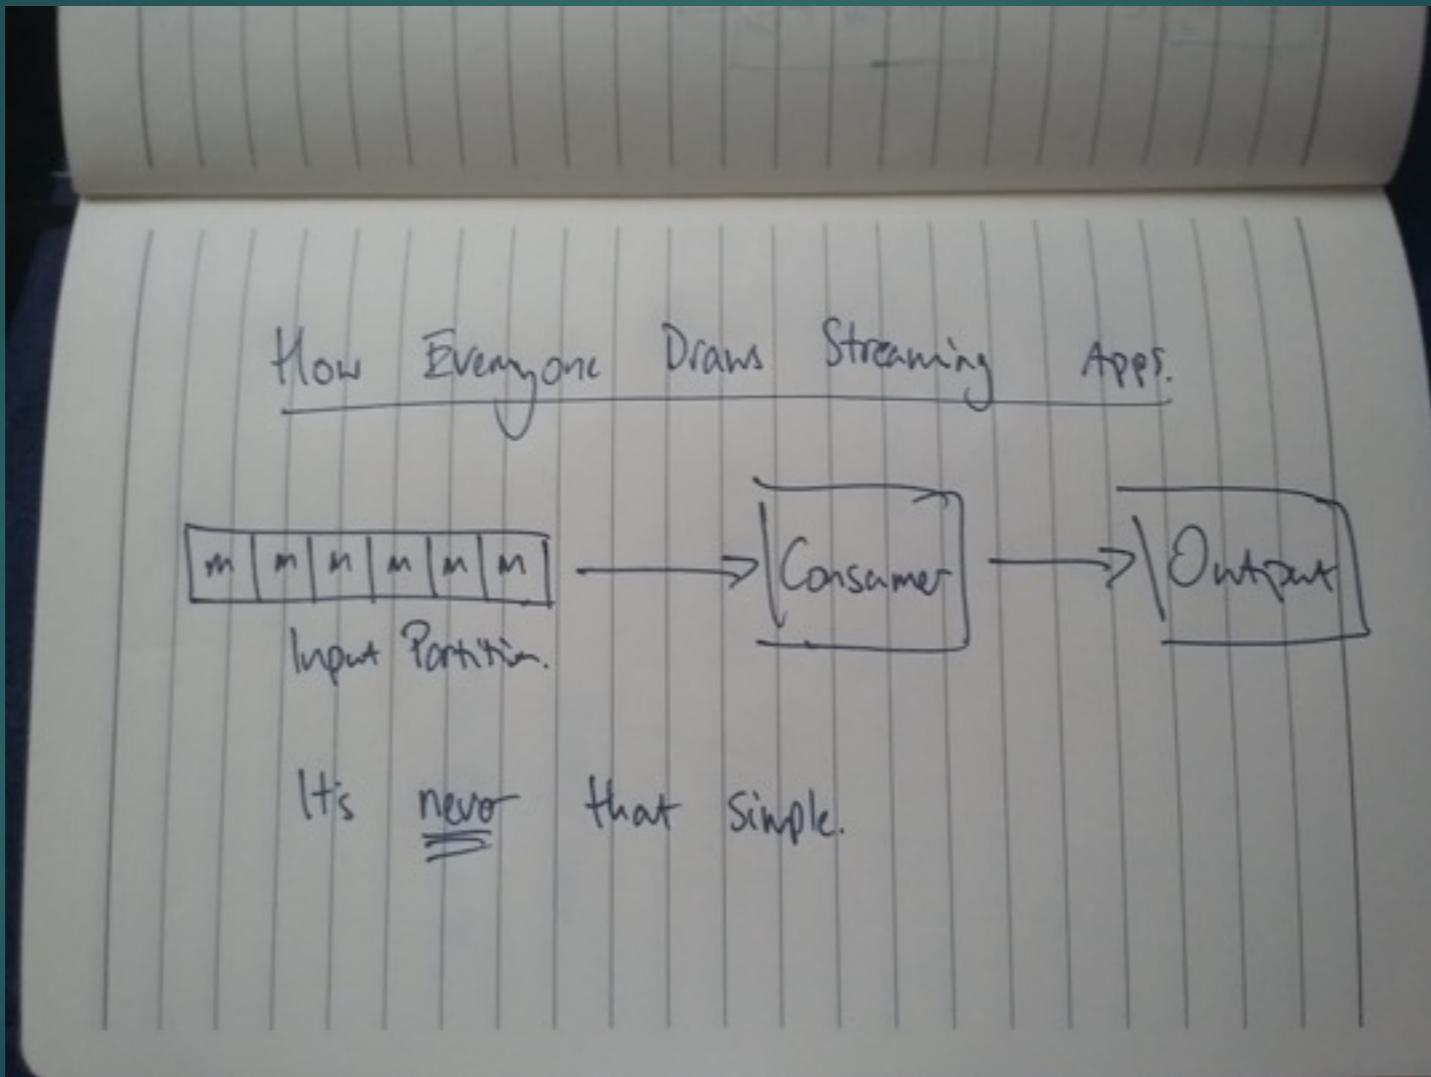


How I Bled All Over
Onyx

One Day I Got a Question

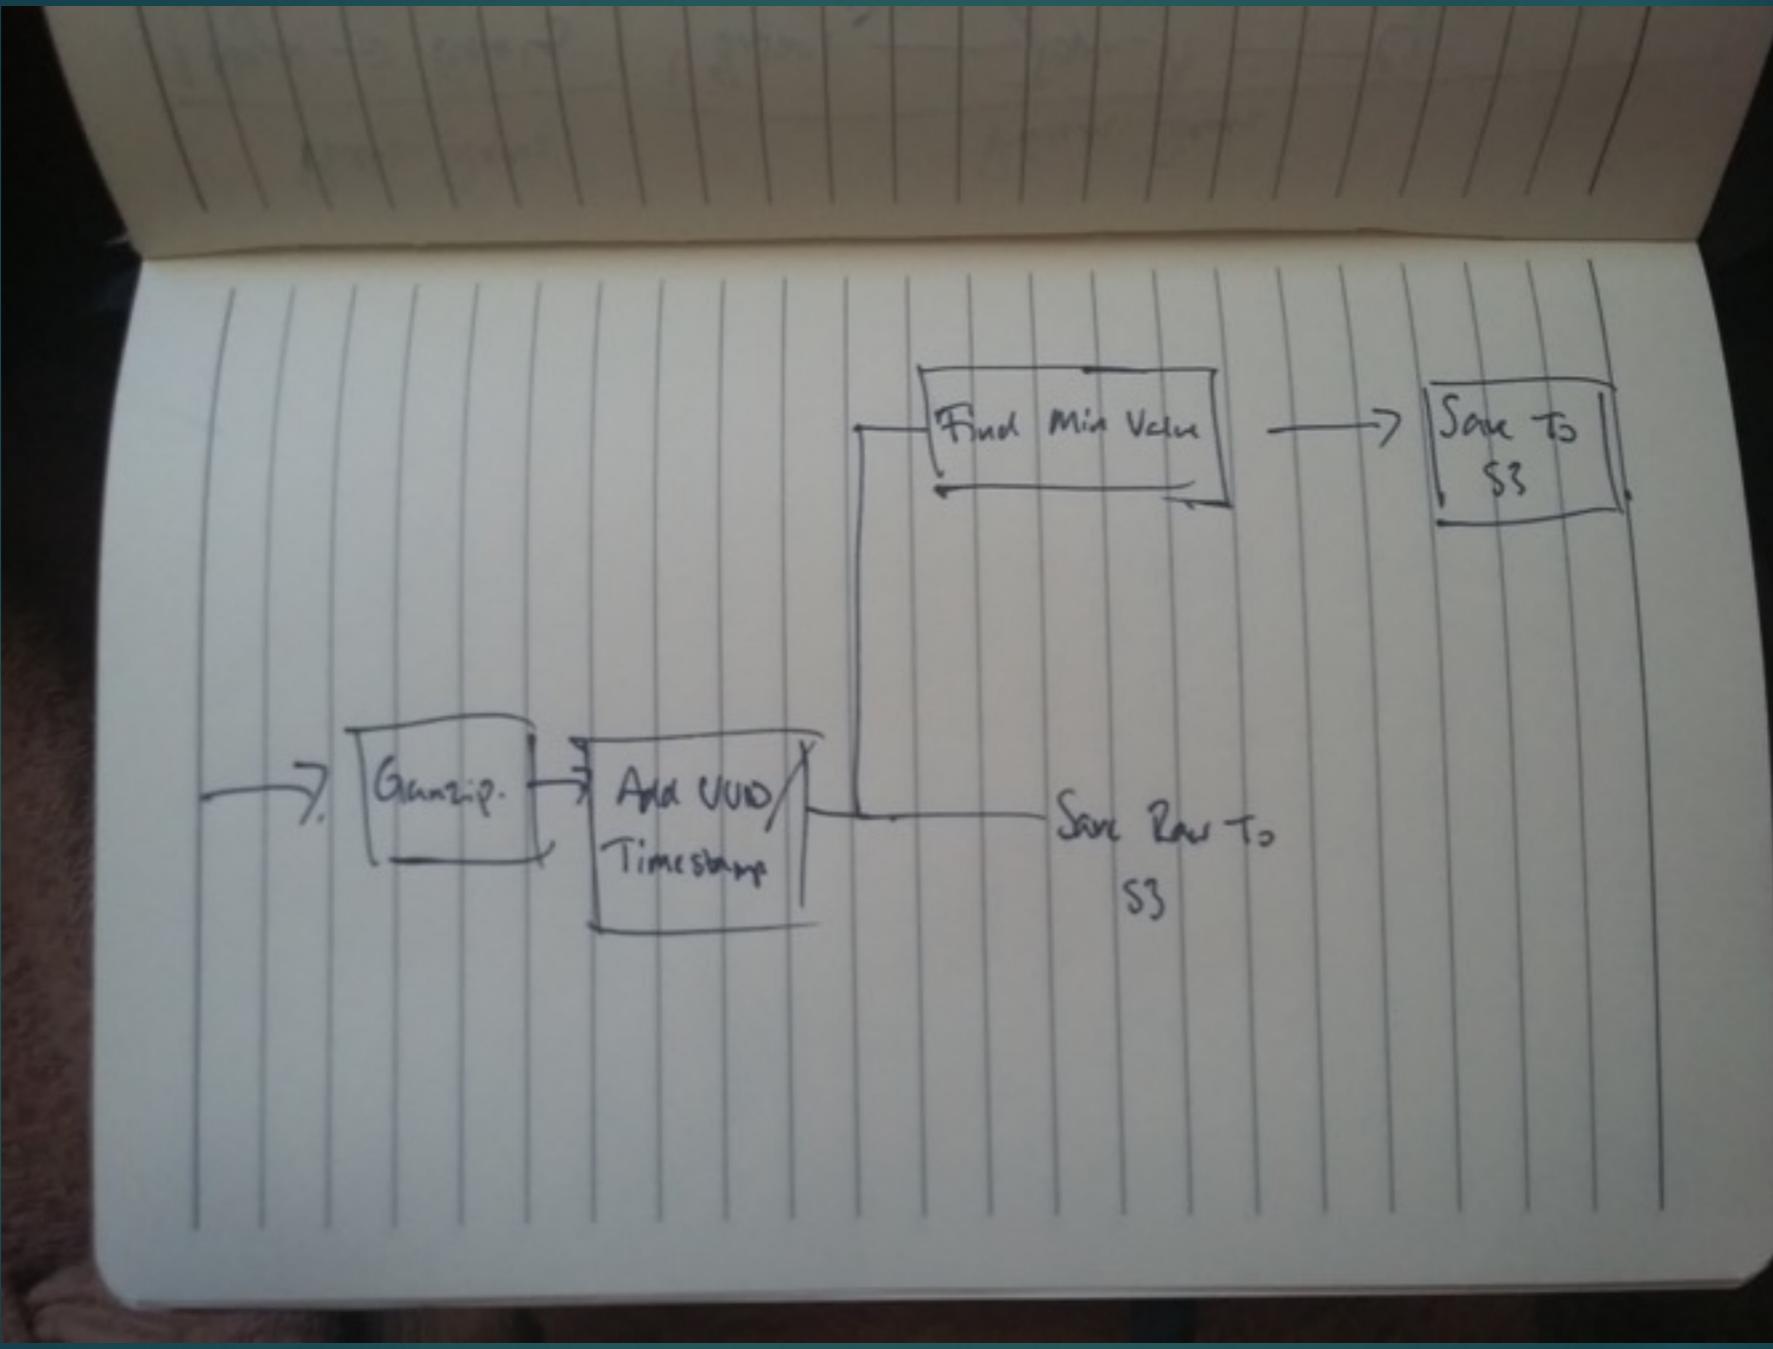
► “We want to stream in 12TB of data a day..... Can you do that?”

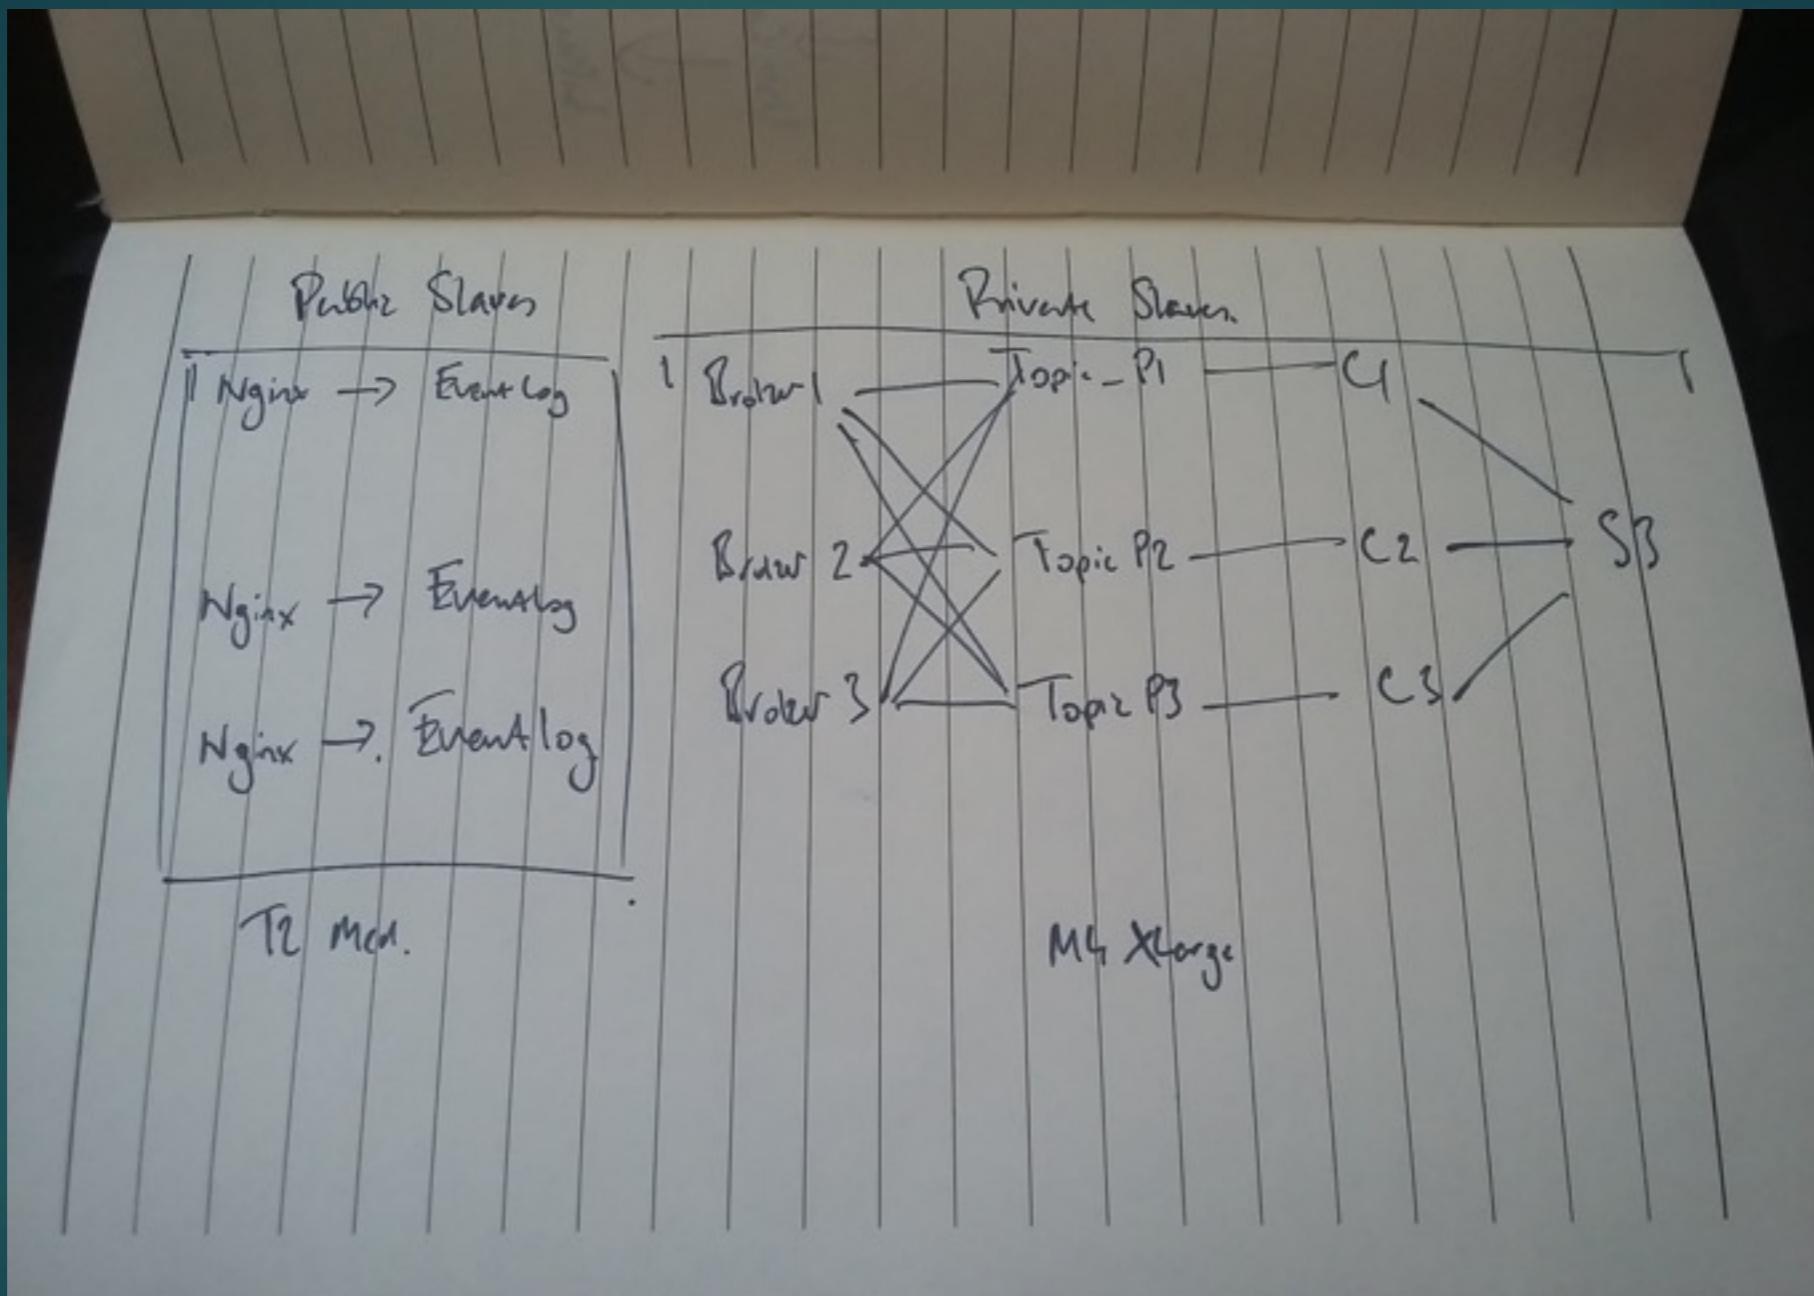
A Streaming Application... kind of.



What is Onyx?

- ▶ Distributed Computation Platform
- ▶ Written 100% in Clojure
- ▶ Great plugin architecture (inputs and outputs)
- ▶ Uses graph Direct Acyclic Graph workflow
- ▶ I spoke about it at ClojureX 2016
- ▶ It's very good....





Phase 1 – Testing at 1% Volume

- ▶ It's working!
- ▶ Stuff's going to S3!
- ▶ Tell the client the good news!

Phase 2 – Testing at 2% Volume

- ▶ It's working!
- ▶ Stuff's going to S3!
- ▶ Tell the client the good news!

Phase 3 – Testing at 5% Volume



Know Thy Framework

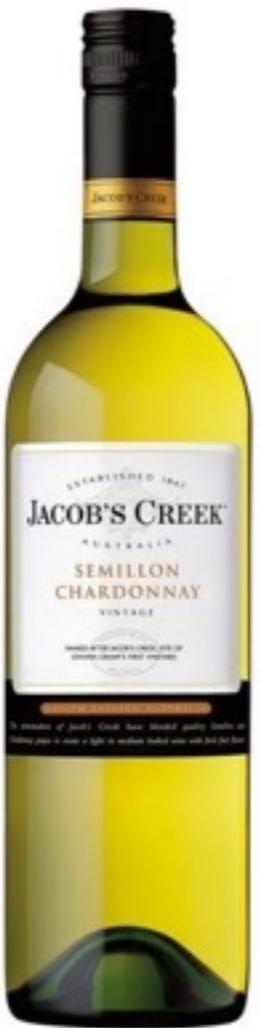
- ▶ Aeron buffer size = $3 \times (\text{batch-size} \times \text{segment size} \times \text{connections})$
 - ▶ $(1 \times 512\text{k} \times 8) \times 3 = 4\text{mb}$
 - ▶ Aeron default buffer is 16mb but then the **twist**
 - ▶ Onyx segment size = `aeron.term.buffer.size / 8`
 - ▶ Max message size of 2mb

Onyx in Docker containers.

- ▶ --shm-size ended up being 10GB
 - ▶ OOM killers were frequent
 - ▶ Java logging is helpful but still hard to deal with.

```
CGROUPS_MEM=$(cat /sys/fs/cgroup/memory/memory.limit_in_bytes)
MEMINFO_MEM=$(($(/proc/meminfo | awk '/MemTotal/ {print $2}') * 1024))
MEM=$(($MEMINFO_MEM > $CGROUPS_MEM ? $CGROUPS_MEM : $MEMINFO_MEM))
JVM_PEER_HEAP_RATIO=${JVM_PEER_HEAP_RATIO:-0.6}
XMX=$(awk '{printf("%d", $1 * $2 / 1024^2)}' <<< "$MEM" \
${JVM_PEER_HEAP_RATIO} "# Use the container memory limit to
set max heap size so that the GC # knows to collect before
it's hard-stopped by the container environment, # causing OOM
exception."
```

One Friday evening...



wine-searcher.com™

2. I did a rewrite in Kafka Streams.

- ▶ Moved the pure Clojure functions into the streams architecture.
- ▶ Used Amazonica to write to AWS S3.

Kafka Streams rewrite...

- ▶ Took 2 hours to write and deploy to DCOS.
- ▶ And, by this stage, was slightly tipsy.

Kafka Streams rewrite...

- ▶ In deployment Kafka Streams didn't touch more than 2% of the CPU/Memory load on the instance.
- ▶ Max memory was 790Mb compared to 8Gb Onyx jobs.

Monday Morning...

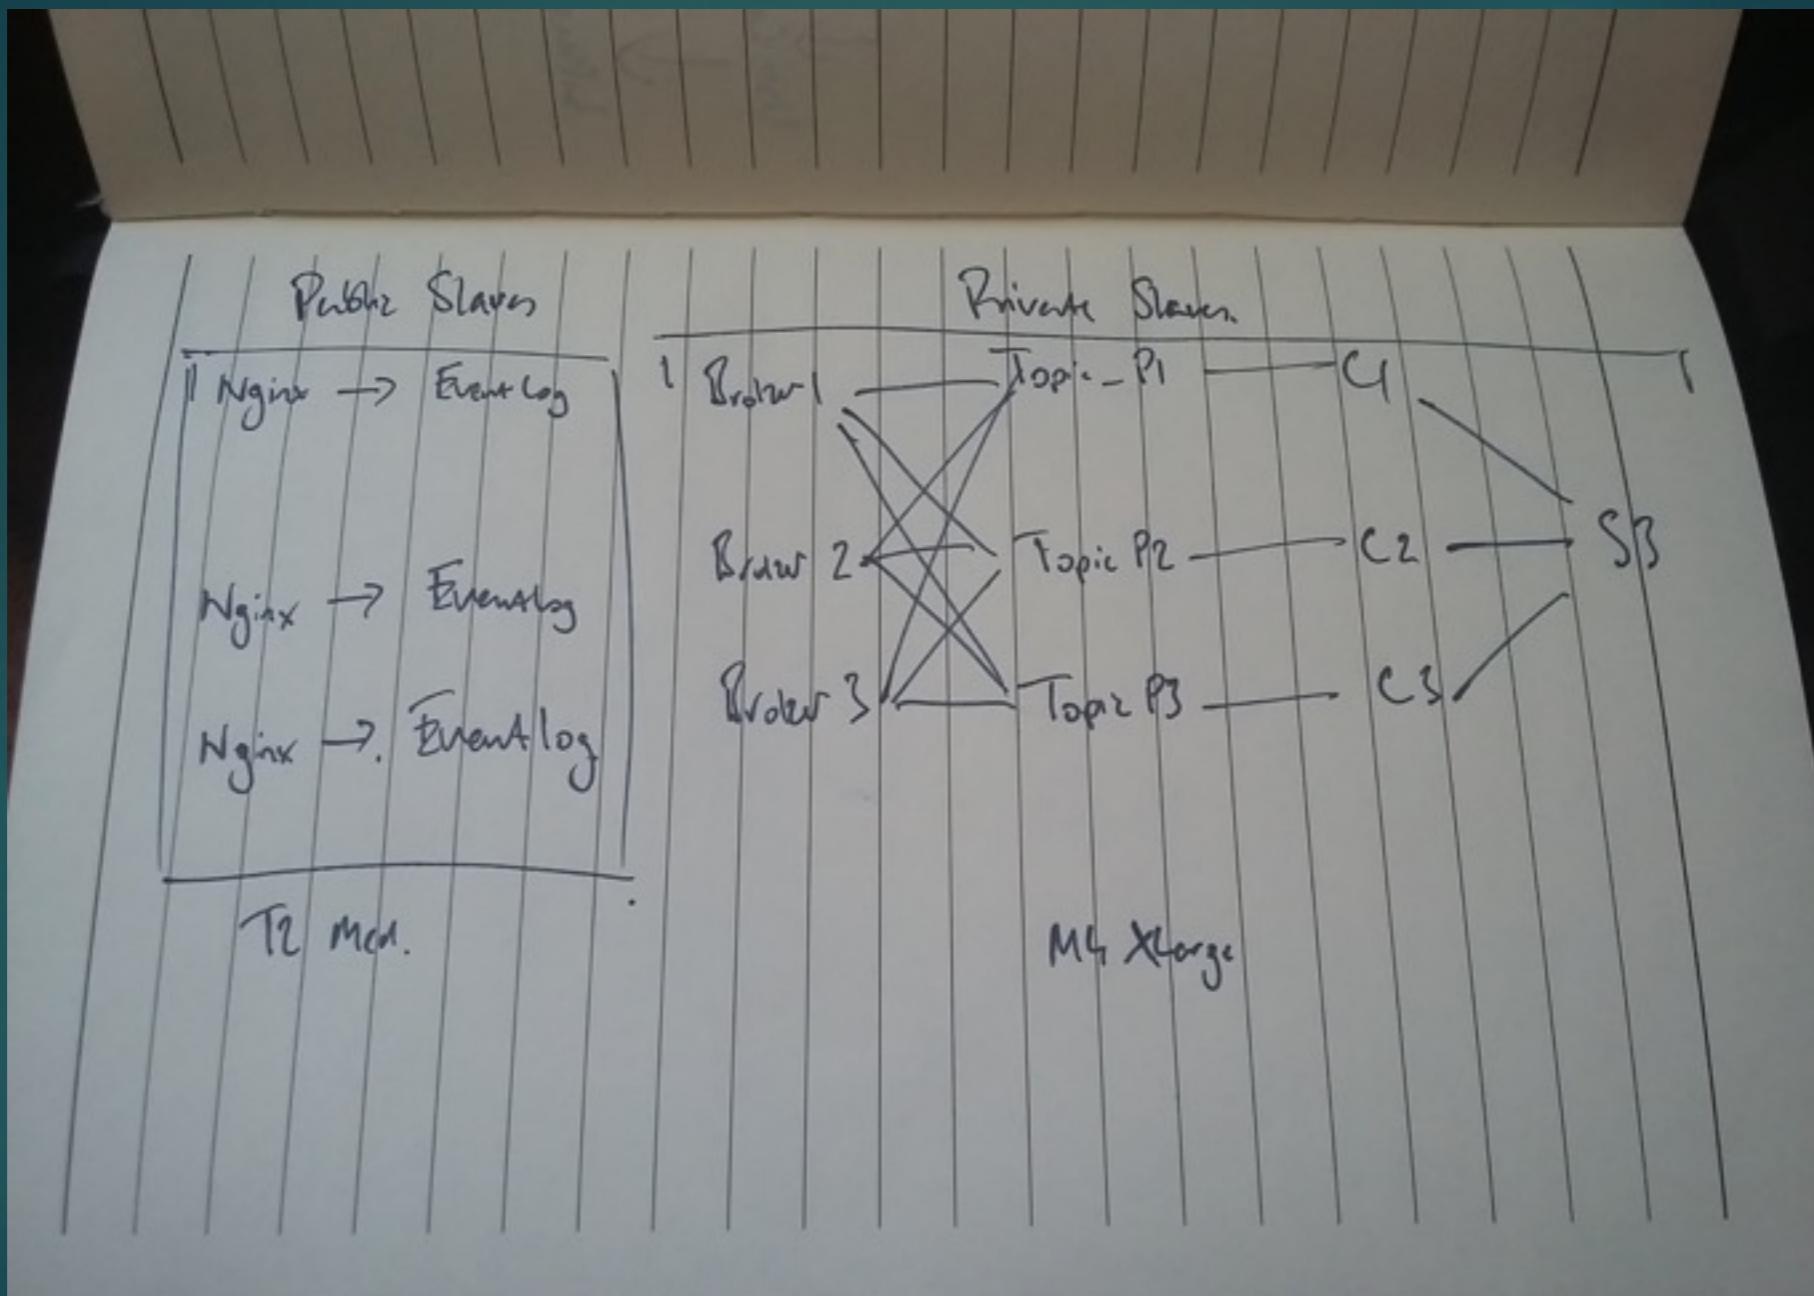
►Now might be a good time to tell
the CTO. ;)

Now lets make it better.

**I need to break down the process
into separate components.**

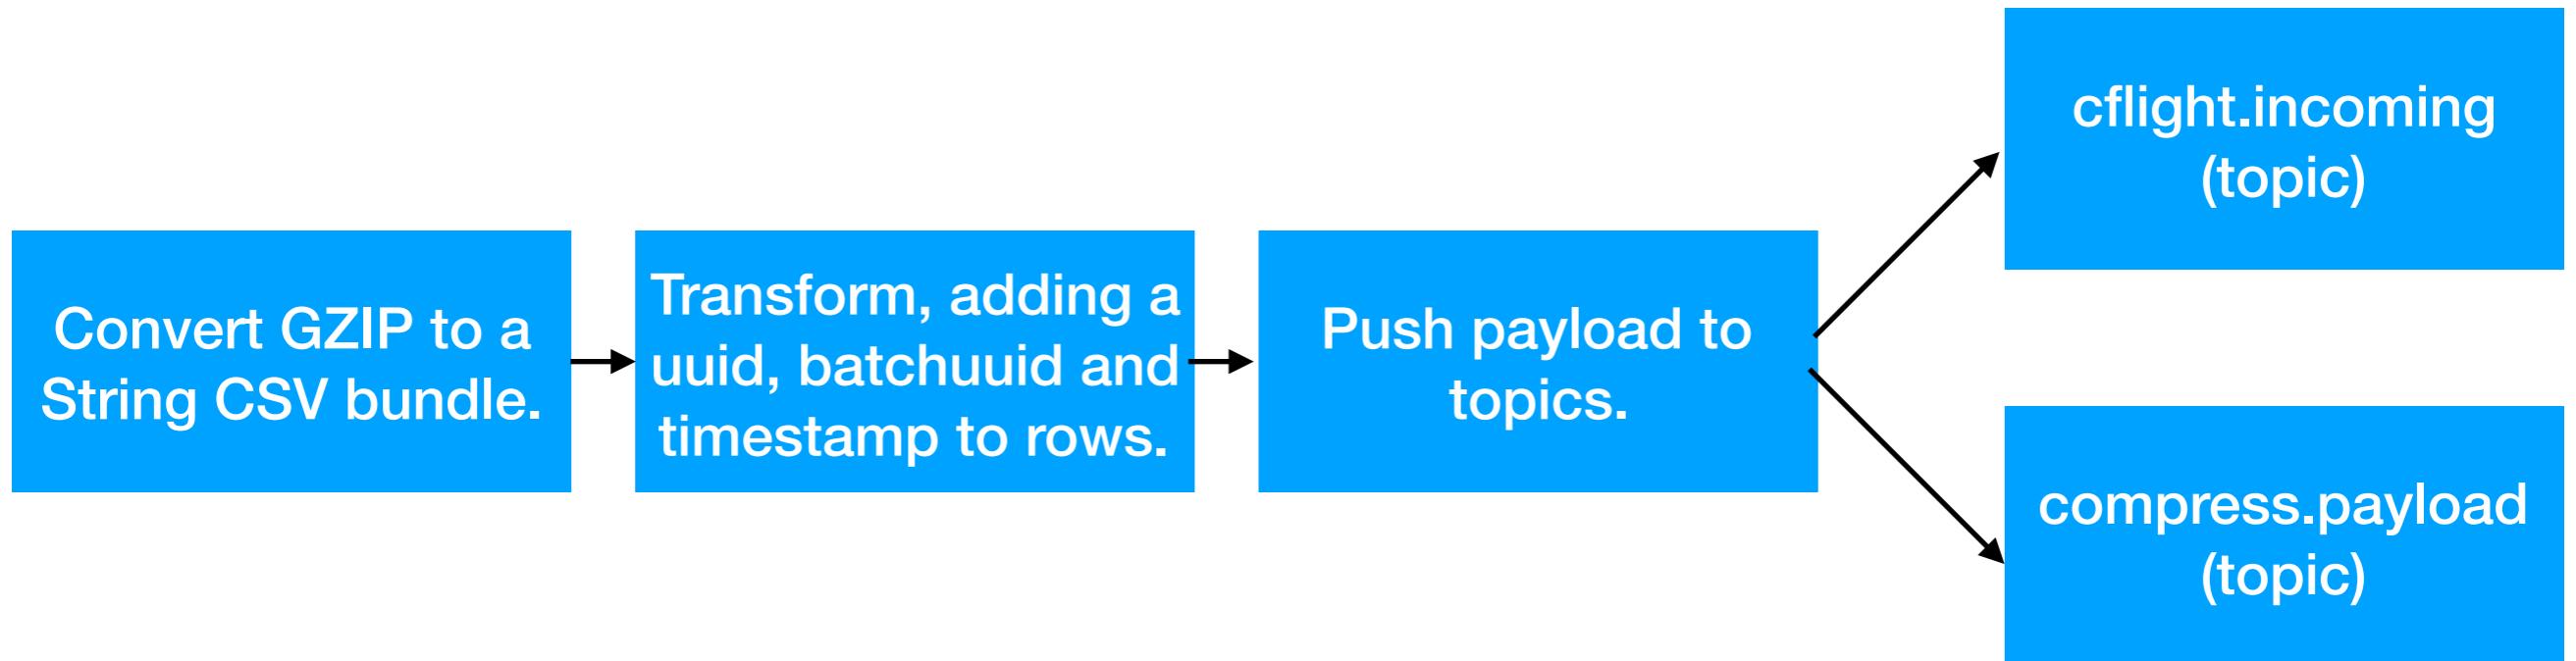
Doing everything in one streaming app is bad!

The Event Log Producer



I'm going to make this into a streaming application.

Streaming Event Log



LHR,IST,2020-08-01,169.99,Y,b6e87fb96f73c119caed08d6e9509c66,39b6aba33172e379c763fcbd1b23aad7,2020-08-01T09:00:00

Deserialise Decompress Step

```
(defn deserialize-gzip-message [bytes]
  (try (-> bytes
              ByteArrayInputStream.
              GZIPInputStream.
              io/reader
              slurp)
       (catch Exception e {:error e})))
```

LHR,IST,2020-08-01,169.99,Y,b6e87fb96f73c119caed08d6e9509c66,39b6aba33172e379c763fcbd1b23aad7,2020-08-01T09:00:00

Transform Step

Define a batch UUID called batchuuid

Define a batch timestamp

Define a StringBuilder outputbundle.

Split bundle by newline:

For each CSV line:

Define a uuid

Append the uuid

Append the batchuuid

Append the timestamp

Append complete line to outputbundle + newline

Push outputbundle string to compress.payload

Push outputbundle string to cflight.incoming

LHR,IST,2020-08-01,169.99,Y,b6e87fb96f73c119caed08d6e9509c66,39b6aba33172e379c763fcbd1b23aad7,2020-08-01T09:00:00

The Cheapest Price Streaming App

Option 1 - Compare each line of the CSV bundle.

```
(defn get-price
  [row]
  (let [first-comma (inc (cstr/index-of row ","))
        last-comma (cstr/index-of row "," first-comma)
        price (subs row first-comma last-comma)]
    (Double/parseDouble price)))

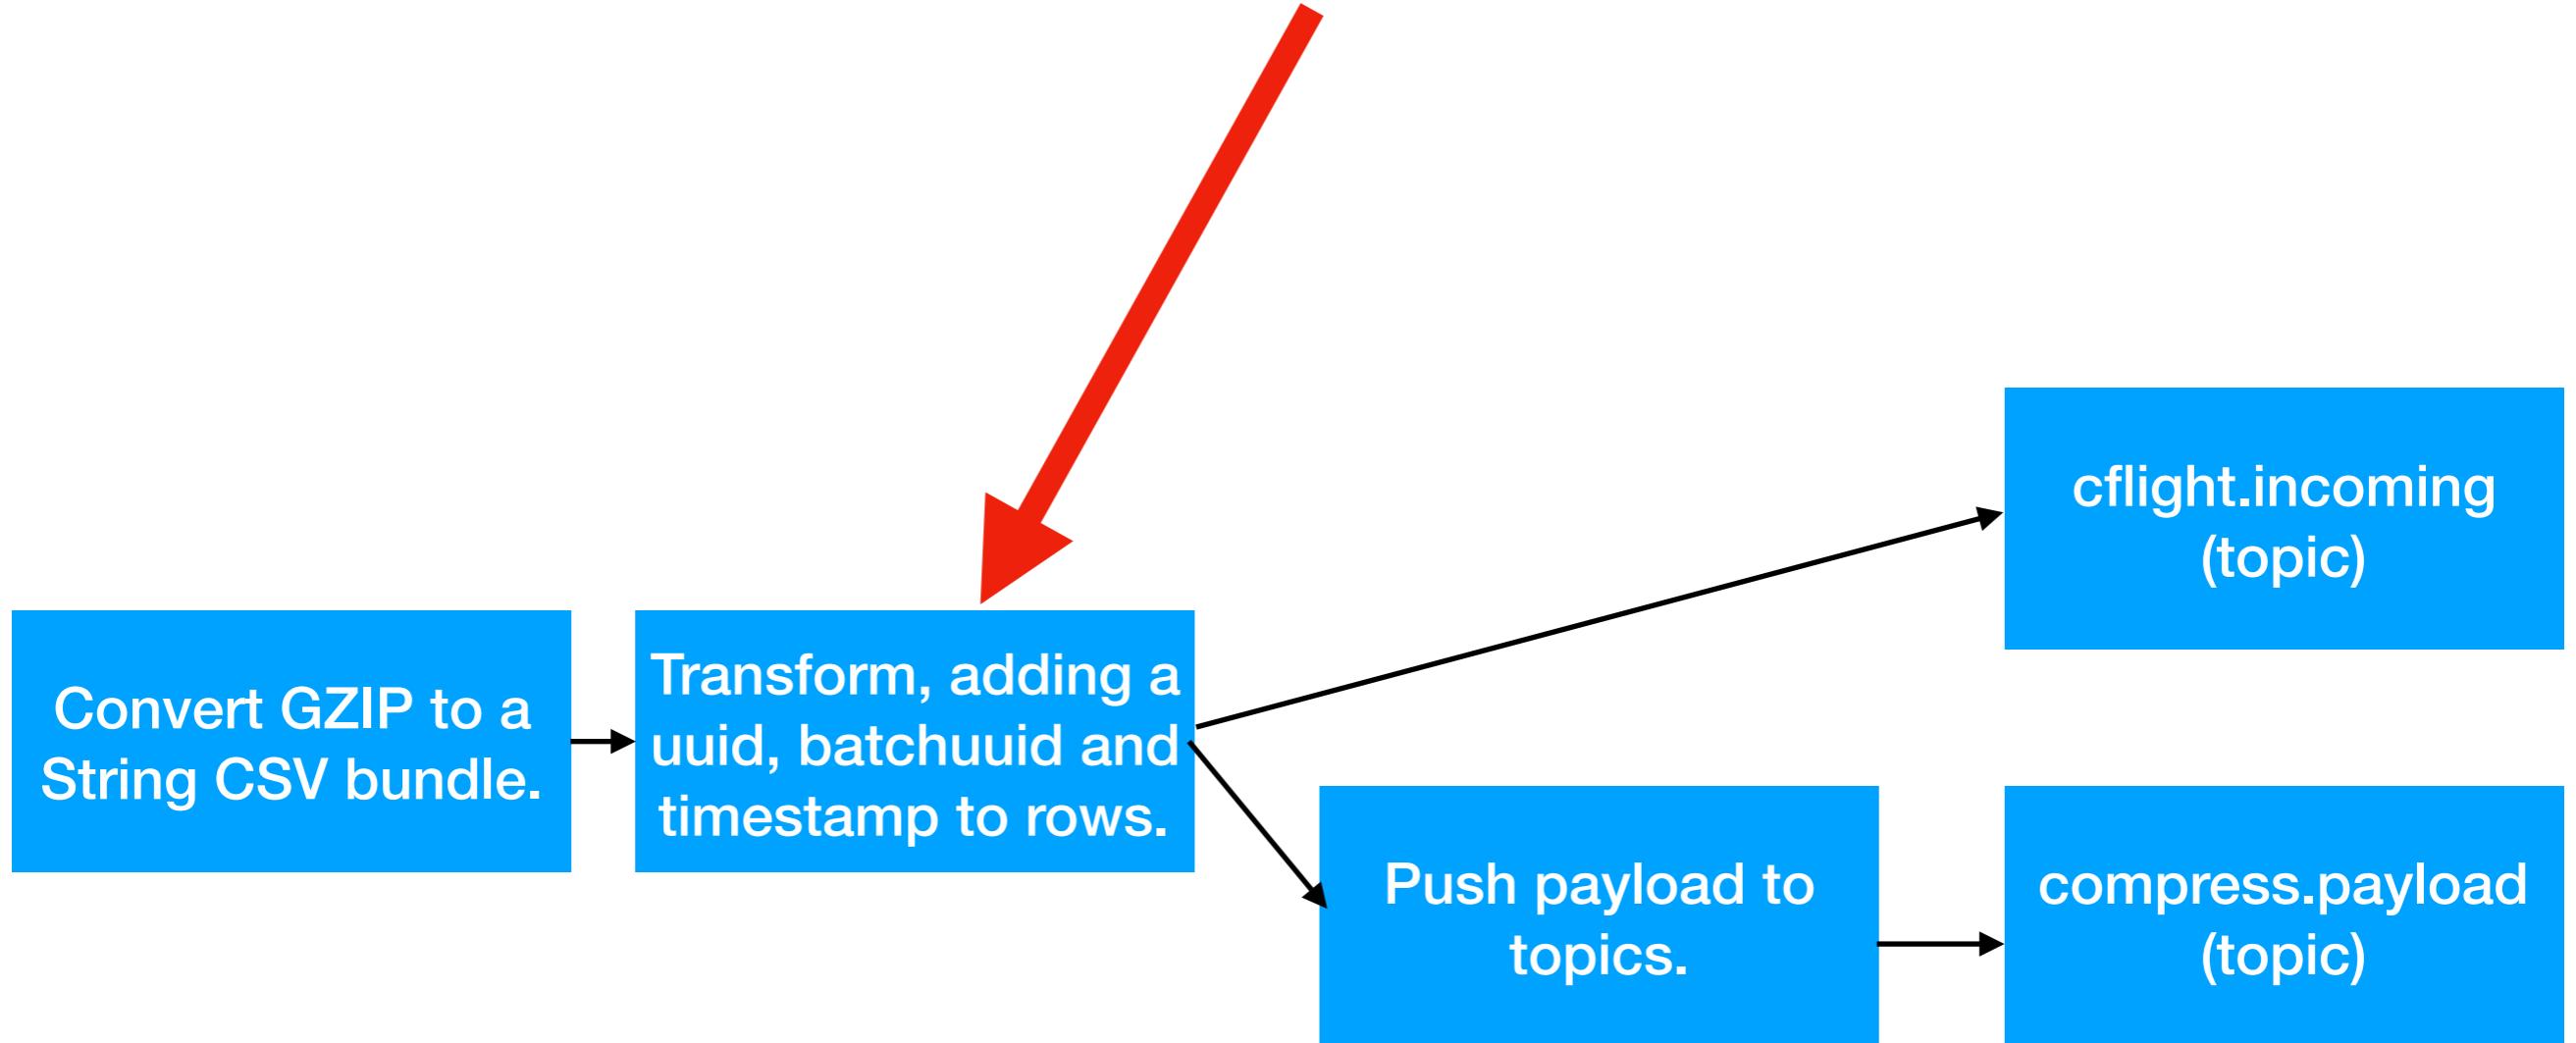
;; Finds the cheapest flight on a row by row comparison.
;; Uses get-price function.
(defn find-cheapest-flight [data]
  (println "Finding cheapest flight")
  (let [rows (cstr/split data #"\n")]
    [price cheapest-flight] (reduce (fn [[cheapest-price cheapest-row :as cheapest]
                                         [candidate-price candidate-row :as candidate]]
                                       (if (> cheapest-price
                                               candidate-price)
                                           candidate
                                           cheapest))
                                       (map (fn [row-o]
                                              [(get-price row-o) row-o])
                                            rows)))
  (str cheapest-flight "\n"))) ;; put the EOL in before we serialise it again
```

OR????????

KSQL?

Option 2 - Step 1: Convert the CSV bundle to AVRO payload

```
{"depiata": "LHR",  
 "arriata": "IST",  
 "searchdate": "2020-08-01",  
 "gbpprice": 169.99,  
 "paxclass": "Y",  
 "uuid": "b6e87fb96f73c119caed08d6e9509c66",  
 "batchuuid": "39b6aba33172e379c763fcbd1b23aad7",  
 "timestamp": "2020-08-01T09:00:00"}
```



LHR,IST,2020-08-01,169.99,Y,b6e87fb96f73c119caed08d6e9509c66,39b6aba33172e379c763fcbd1b23aad7,2020-08-01T09:00:00

Event Log Transform Step Amended

Define a batch UUID called batchuuid

Define a batch timestamp

Define a StringBuilder outputbundle.

Split bundle by newline:

For each CSV line:

Define a uuid

Append the uuid

Append the batchuuid

Append the timestamp

Append complete line to outputbundle + newline

Convert line to AVRO

Push AVRO payload to cflight.incoming topic

Push outputbundle string to compress.payload

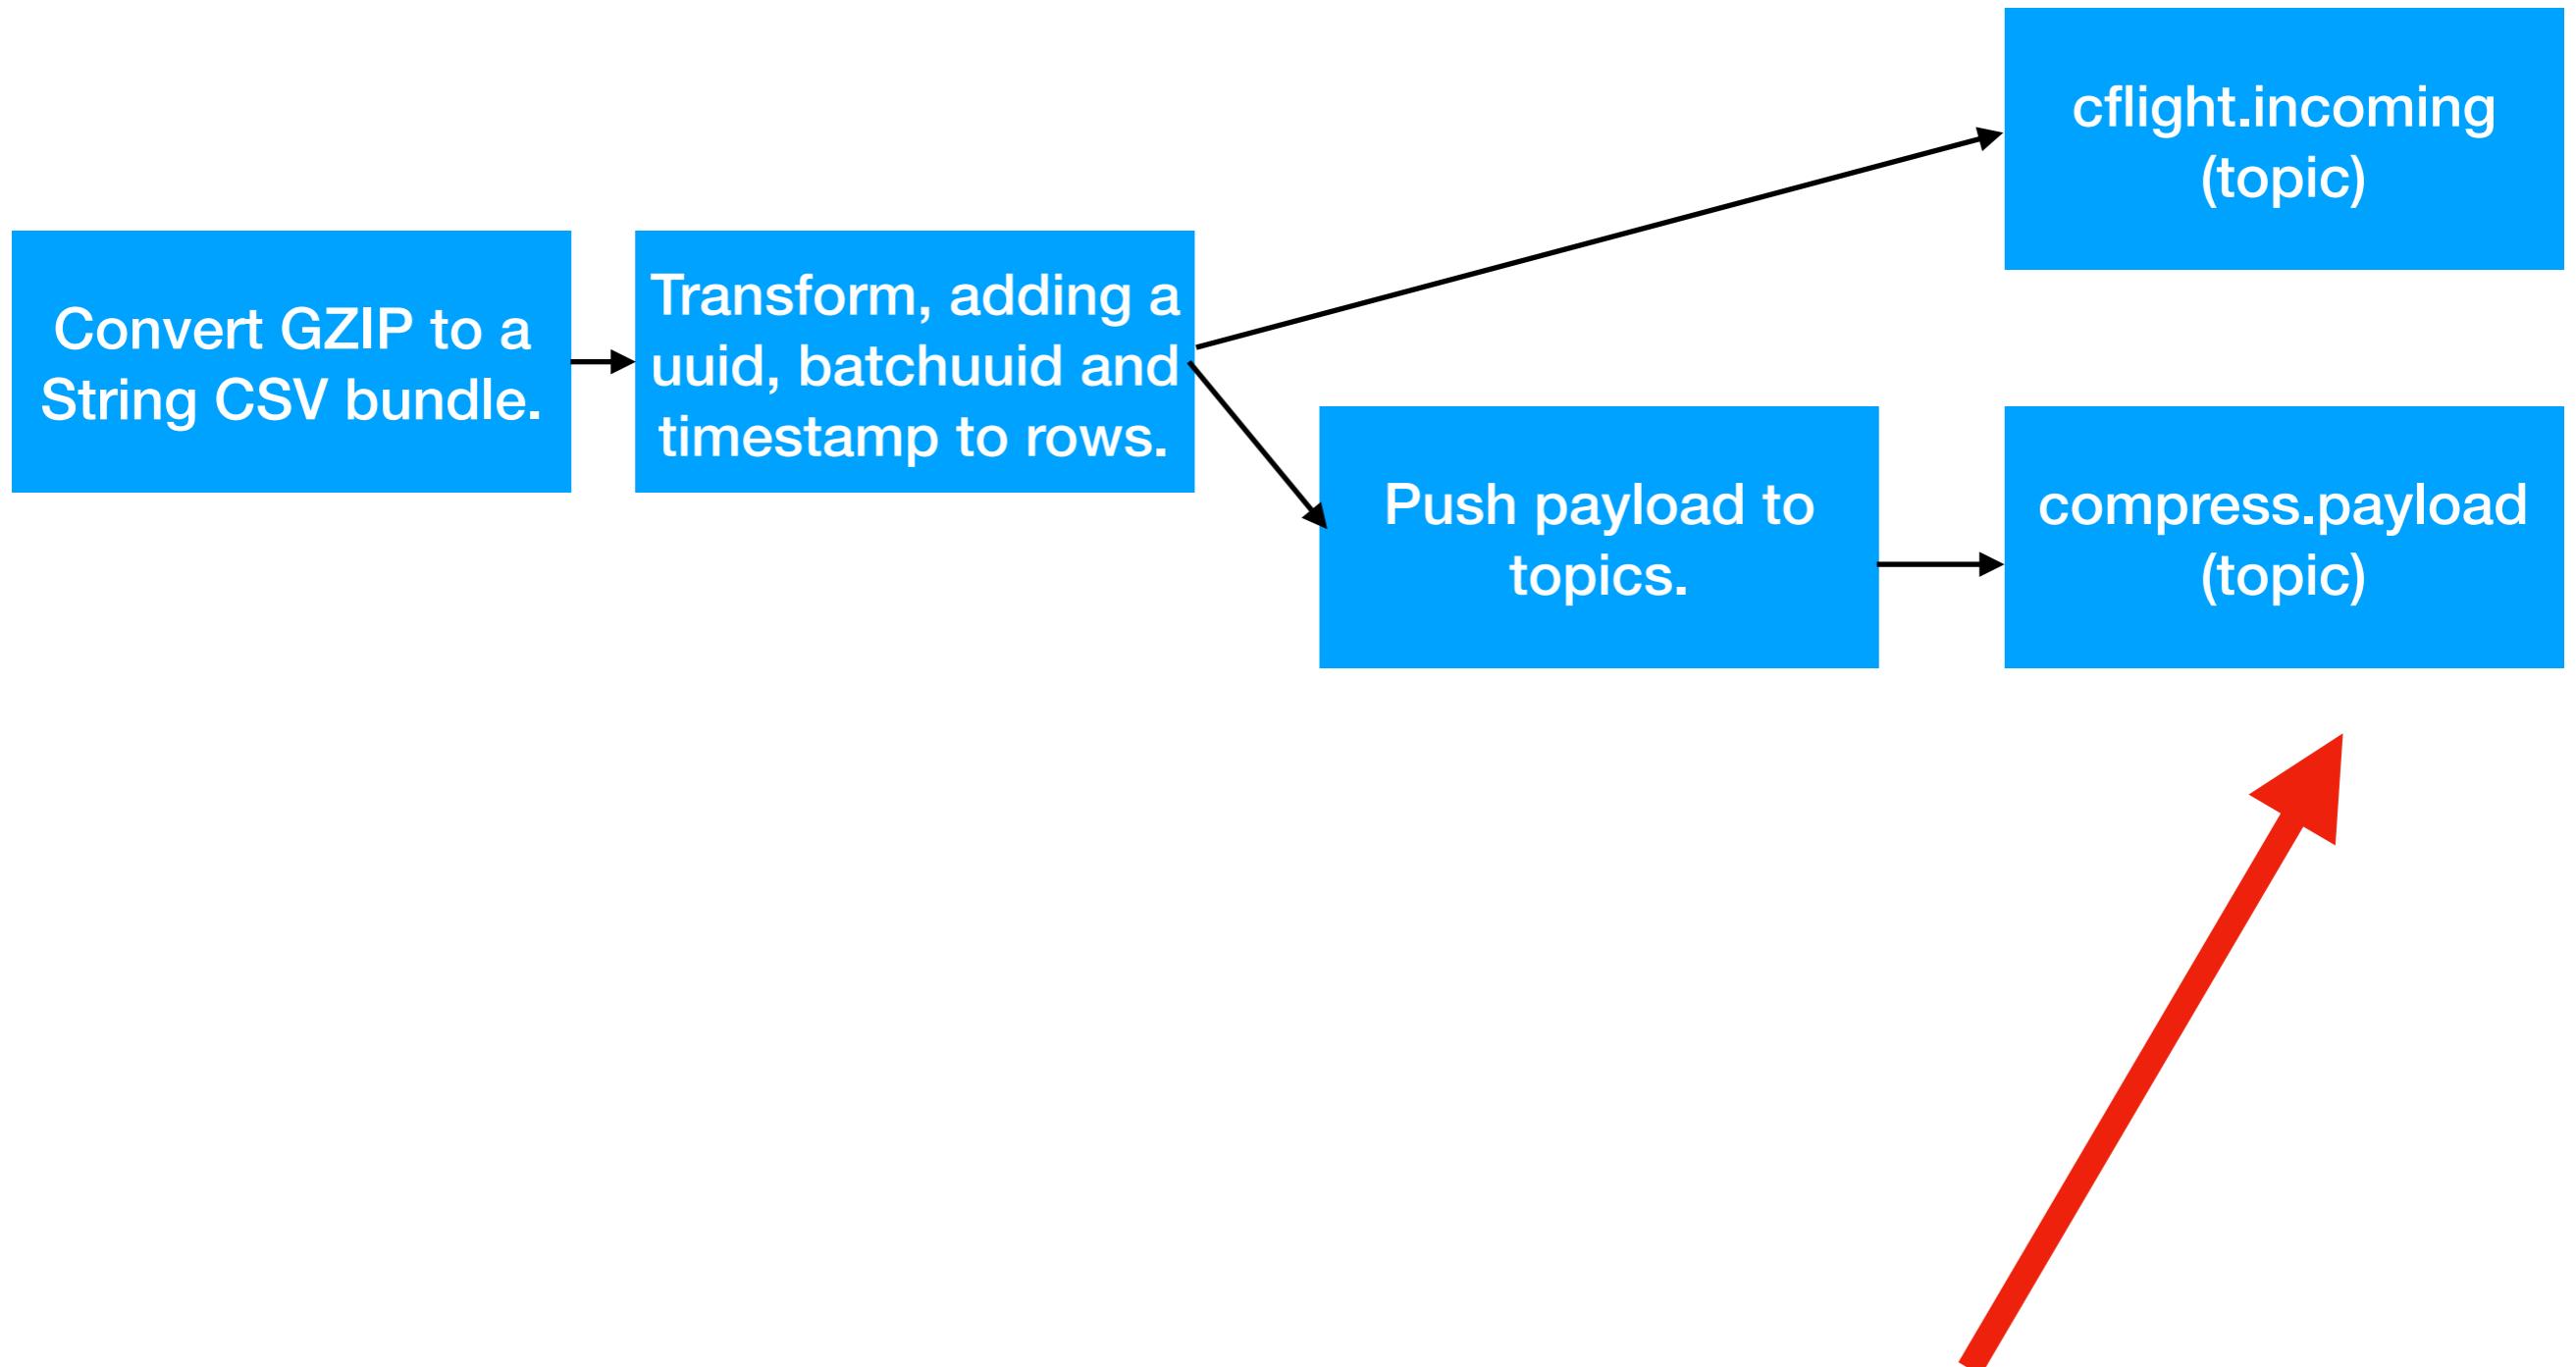
LHR,IST,2020-08-01,169.99,Y,b6e87fb96f73c119caed08d6e9509c66,39b6aba33172e379c763fcbd1b23aad7,2020-08-01T09:00:00

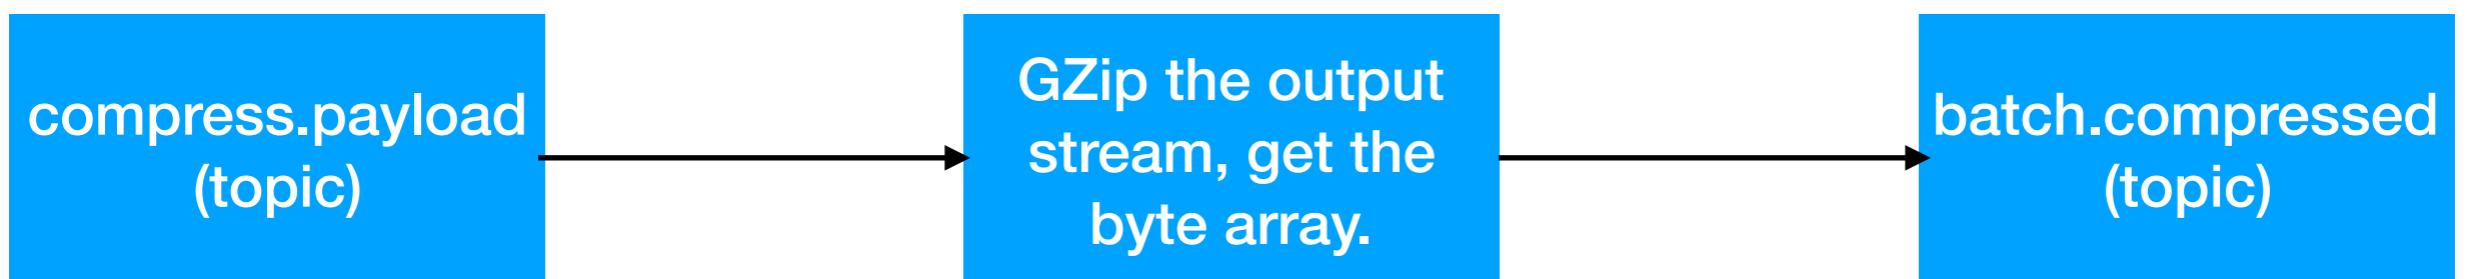
Option 2 - Step 2: Craft a KSQL Job

```
CREATE STREAM CFLIGHT (.....)
    WITH (KAFKA_TOPIC='cflight.incoming',
          PARTITIONS=1,
          VALUE_FORMAT='avro');
```

```
CREATE STREAM CFLIGHT_CHEAPEST AS
    SELECT UUID, BATCHUUID,
           MIN(GBPPRICE) AS CHEAPEST_QUOTE
    FROM CFLIGHT
    GROUP BY BATCHUUID
    EMIT CHANGES;
```

The Batch Compressor Streaming App





Compress and Serialise Step

```
(def gzip-serializer-fn
  (fn [output-str]
    (let [out (java.io.ByteArrayOutputStream.)]
      (do (doto (java.io.BufferedOutputStream.
                  (java.util.zip.GZIPOutputStream. out))
            (.write (.getBytes output-str)))
          (.close)))
        (.toByteArray out)))))
```

Persist to S3

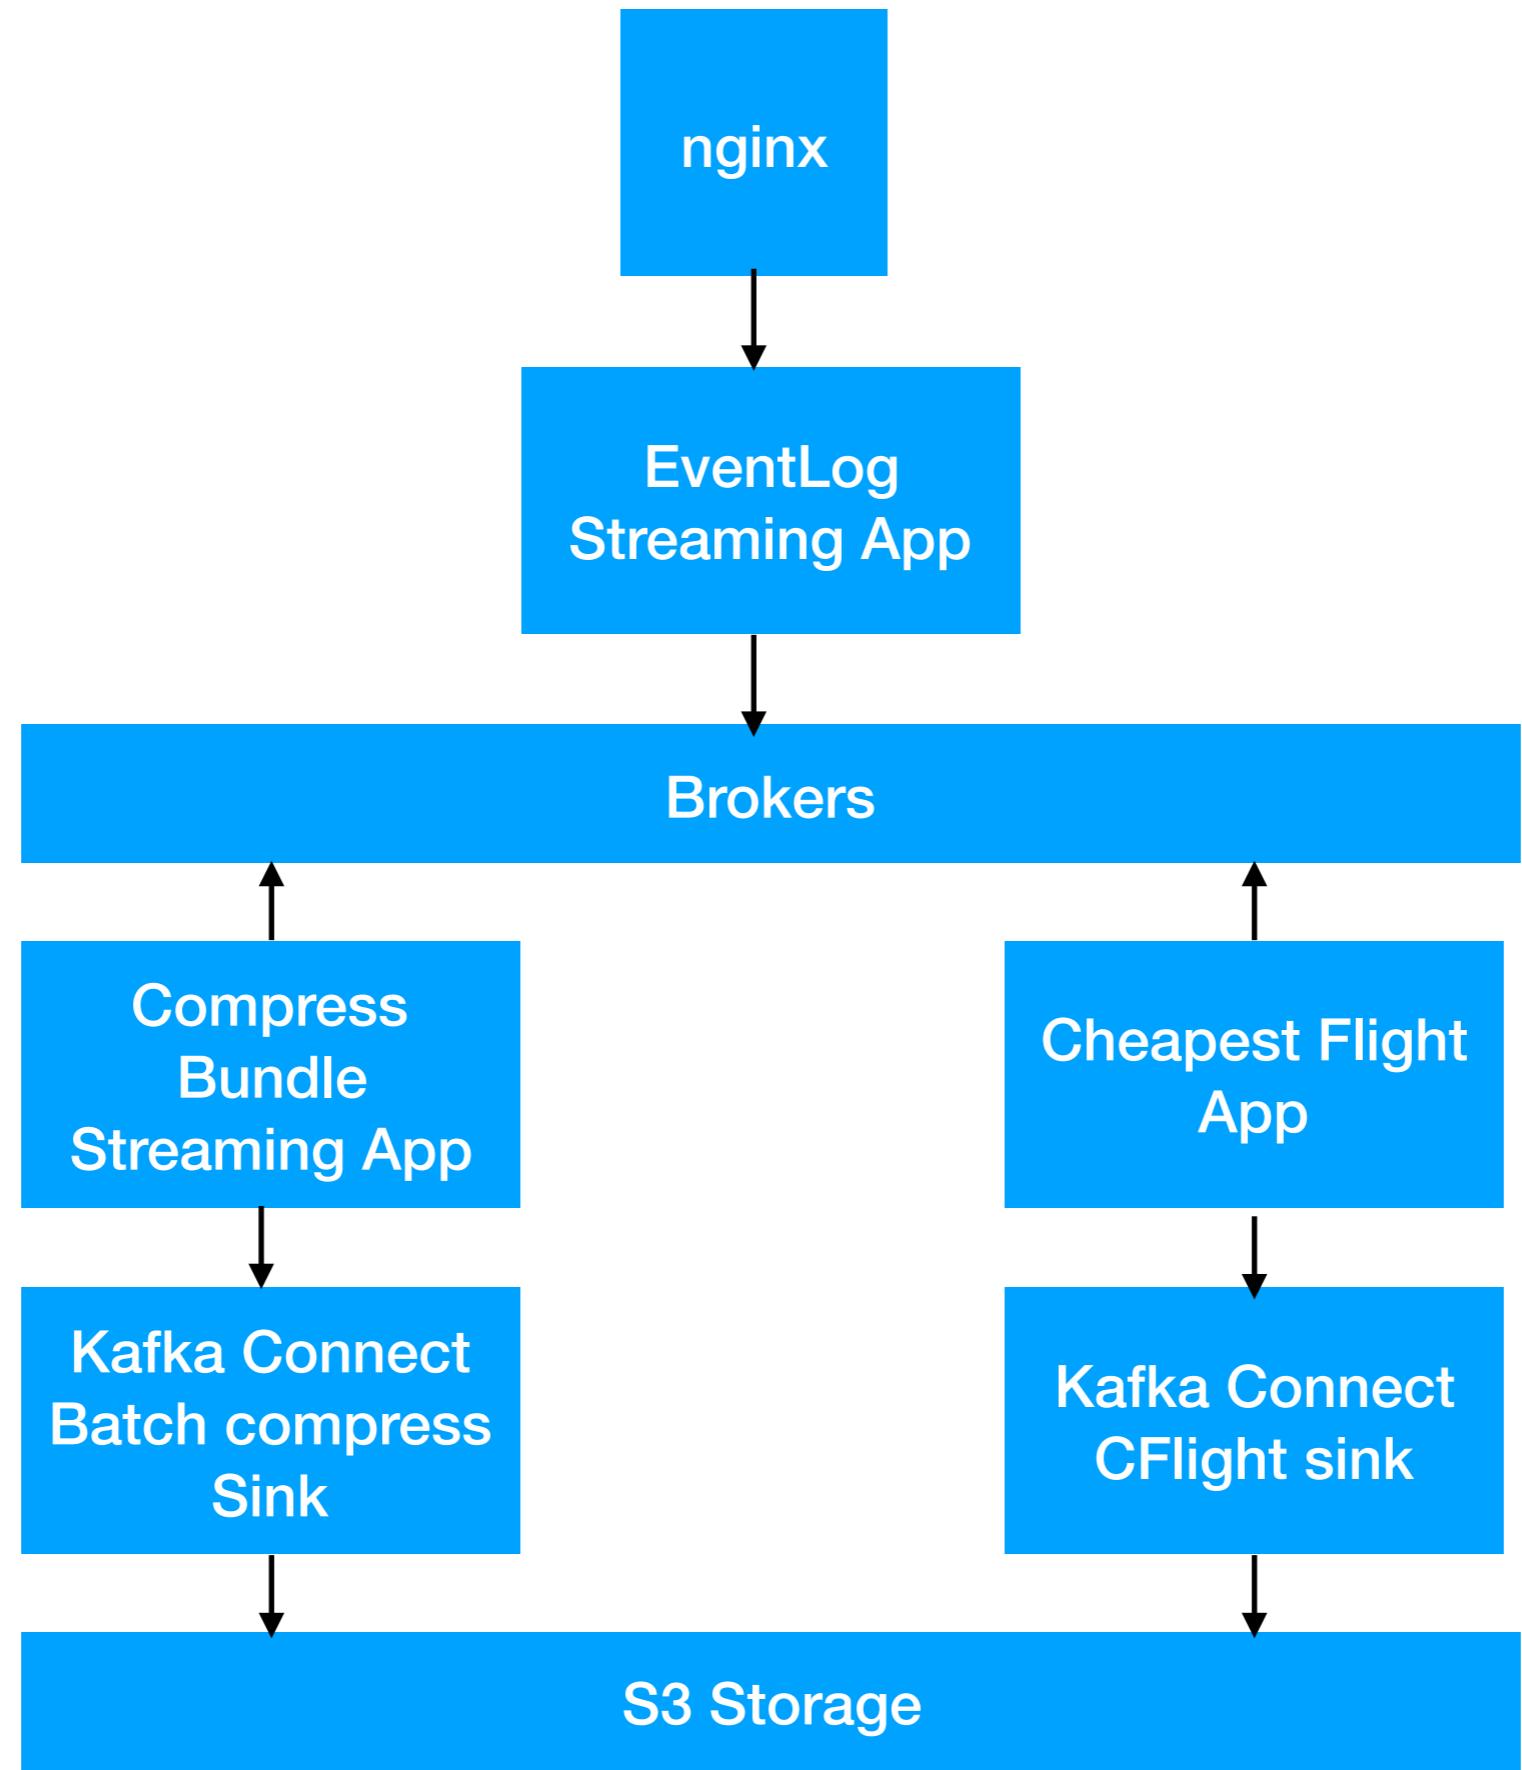
Kafka Connect

**Use Prebundled connectors
where you can.**

**Also means you have access to further filter
and transform steps if required.**

```
{"name": "cflight.sink",
"config":
{"connector.class": "io.confluent.connect.s3.S3SinkConnector",
"tasks.max": "1",
"topics": "CHEAPEST_FLIGHT",
"s3.region": "us-east-1",
"s3.bucket.name": "cflight.bucket",
"s3.part.size": "5242880",
"flush.size": "1",
"storage.class": "io.confluent.connect.s3.storage.S3Storage",
"format.class": "io.confluent.connect.s3.format.avro.AvroFormat",
"partitioner.class": "io.confluent.connect.storage.partition.DefaultPartitioner",
"schema.compatibility": "NONE"}}
```

```
{"name": "batchcompress.sink",
"config":
{"connector.class": "io.confluent.connect.s3.S3SinkConnector",
"tasks.max": "1",
"topics": "batch.compressed",
"s3.region": "us-east-1",
"s3.bucket.name": "flightbatch.bucket",
"s3.part.size": "5242880",
"flush.size": "1",
"storage.class": "io.confluent.connect.s3.storage.S3Storage",
"format.class": "io.confluent.connect.s3.format.ByteArrayFormat",
"partitioner.class": "io.confluent.connect.storage.partition.DefaultPartitioner",
"schema.compatibility": "NONE"}}
```



Thank you.

Many thanks to Shay and David for organising, everyone who attended and sent kind wishes. Lastly, a huge thank you to MeetupCat.



Photo supplied by @jbfletch_